

A PYTHON PROGRAM TO IMPLEMENT ADA BOOSTING

Expt no. 8A

Name: Jayasuriya.j

Roll no: 241801102

PROGRAM:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeClassifier,
plot_tree
from mlxtend.plotting import
plot_decision_regions

df = pd.DataFrame({
    'X1':[1,2,3,4,5,6,6,7,9,9],
    'X2':[5,3,6,8,1,9,5,8,9,2],
    'label':[1,1,0,1,0,1,0,1,0,0]
})
print("Original Data:\n", df)
```

```
# Scatter plot
plt.figure(figsize=(6,4))
plt.scatter(df['X1'], df['X2'], c=df['label'],
cmap='coolwarm', edgecolors='black')
plt.xlabel("X1"); plt.ylabel("X2")
plt.title("Scatter Plot of Data")
plt.show()
```

```
df['weights'] = 1/df.shape[0]
print("\nInitial Weights:\n", df)

x = df[['X1','X2']].values
y = df['label'].values

dt1 = DecisionTreeClassifier(max_depth=1)
dt1.fit(x, y)

plt.figure(figsize=(6,4))
plot_tree(dt1, filled=True,
          feature_names=['X1','X2'], class_names=['0','1'])
plt.title("Decision Stump dt1")
plt.show() plot_decision_regions(x, y,
                                 clf=dt1, legend=2) plt.title("Decision Region
dt1")
plt.show()

df['y_pred'] = dt1.predict(x)
print("\nAfter dt1 Predictions:\n",
      df[['X1','X2','label','weights','y_pred']])

def calculate_alpha(error):
    return 0.5 * np.log((1-error)/error)

error1 = (df['label'] != df['y_pred']).mean()
alpha1 = calculate_alpha(error1)
print("\nError1 =", error1)
print("Alpha1 =", alpha1)
```

```

def update_row_weights(row, alpha):
    if row['label'] == row['y_pred']:
        return row['weights'] * np.exp(-alpha)
    else:
        return row['weights'] * np.exp(alpha)

df['updated_weights'] = df.apply(lambda row:
    update_row_weights(row, alpha1), axis=1)
print("\nUpdated Weights (Round 1):\n",
df[['X1','X2','label','weights','y_pred','updated_wei
ghts']])

```



```

# Normalized weights
df['normalized_weights'] =
df['updated_weights'] /
df['updated_weights'].sum()
print("\nNormalized Weights (Round 1):\n",
df[['X1','X2','label','normalized_weights']])
df['cumsum_upper'] =
np.cumsum(df['normalized_weights'])
df['cumsum_lower'] =
df['cumsum_upper'] -
df['normalized_weights']
print("\nCumsum Bounds (Round
1):\n",
df[['cumsum_lower','cumsum_upper']])

```



```


indices = []
for i in range(df.shape[0]):
    r = np.random.random()

```

```
for idx, row in df.iterrows():
    if row['cumsum_lower'] <= r <=
row['cumsum_upper']:
    indices.append(idx)
    break


second_df = df.iloc[indices].copy()
second_df['weights'] = 1/df.shape[0]
print("\nSecond Dataset (Weighted Sampling):\n",
second_df[['X1','X2','label','weights']])
x2 = second_df[['X1','X2']].values
y2 = second_df['label'].values
dt2 = DecisionTreeClassifier(max_depth=1)
dt2.fit(x2, y2)
plt.figure(figsize=(6,4))
plot_tree(dt2, filled=True,
feature_names=['X1','X2'], class_names=['0','1'])
plt.title("Decision Stump dt2")
plt.show()
plot_decision_regions(x2, y2, clf=dt2, legend=2)
plt.title("Decision Region dt2")
plt.show()
```

OUTPUT:



Initial Weights:

	x1	x2	label	weights
0	1	5	1	0.1
1	2	3	1	0.1
2	3	6	0	0.1
3	4	8	1	0.1
4	5	1	0	0.1
5	6	9	1	0.1
6	6	5	0	0.1
7	7	8	1	0.1
8	9	9	0	0.1
9	9	2	0	0.1



```

After dt1 Predictions:
  X1  X2  label  weights  y_pred
0   1   5     1     0.1      1
1   2   3     1     0.1      1
2   3   6     0     0.1      0
3   4   8     1     0.1      0
4   5   1     0     0.1      0
5   6   9     1     0.1      0
6   6   5     0     0.1      0
7   7   8     1     0.1      0
8   9   9     0     0.1      0
9   9   2     0     0.1      0

Error1 = 0.3
Alpha1 = 0.42364893019360184

Updated Weights (Round 1):
  X1  X2  label  weights  y_pred  updated_weights
0   1   5     1     0.1      1      0.065465
1   2   3     1     0.1      1      0.065465
2   3   6     0     0.1      0      0.065465
3   4   8     1     0.1      0      0.152753
4   5   1     0     0.1      0      0.065465
5   6   9     1     0.1      0      0.152753
6   6   5     0     0.1      0      0.065465
7   7   8     1     0.1      0      0.152753
8   9   9     0     0.1      0      0.065465
9   9   2     0     0.1      0      0.065465
  
```


```

Normalized Weights (Round 1):
  x1  x2  label  normalized_weights
0   1   5       1      0.071429
1   2   3       1      0.071429
2   3   6       0      0.071429
3   4   8       1      0.166667
4   5   1       0      0.071429
5   6   9       1      0.166667
6   6   5       0      0.071429
7   7   8       1      0.166667
8   9   9       0      0.071429
9   9   2       0      0.071429

Cumsum Bounds (Round 1):
  cumsum_lower  cumsum_upper
0   0.000000    0.071429
1   0.071429    0.142857
2   0.142857    0.214286
3   0.214286    0.380952
4   0.380952    0.452381
5   0.452381    0.619048
6   0.619048    0.690476
7   0.690476    0.857143
8   0.857143    0.928571
9   0.928571    1.000000

Second Dataset (Weighted Sampling):
  x1  x2  label  weights
9   9   2       0      0.1
3   4   8       1      0.1
6   6   5       0      0.1
7   7   8       1      0.1
9   9   2       0      0.1
3   4   8       1      0.1
5   6   9       1      0.1
7   7   8       1      0.1
2   3   6       0      0.1

```



A PYTHON PROGRAM TO IMPLEMENT GRADIENT BOOSTING

Expt No. 8B

PROGRAM:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeRegressor

plt.style.use("seaborn-v0_8")
np.random.seed(42)

X = np.random.rand(100, 1) - 0.5
y = 3 * (X[:, 0] ** 2) + 0.05 * np.random.randn(100)

df = pd.DataFrame({"X": X.reshape(-1), "y": y})

plt.figure(figsize=(6,4))
plt.scatter(df["X"], df["y"], s=25)
plt.title("X vs y")
plt.xlabel("X"); plt.ylabel("y")
plt.show()

def gradient_boost(X, y, n_estimators=5, lr=1.0, max_leaf_nodes=8,
random_state=42): rng = np.random.RandomState(random_state)

    X = np.asarray(X).reshape(-1, 1)
    y = np.asarray(y)

    F = np.full_like(y, y.mean(), dtype=float)
    trees = []
    xg = np.linspace(X.min() - 0.05, X.max() + 0.05, 600).reshape(-1, 1)
    curves = []
    stage_metrics = []

    for i in range(n_estimators):
        tree = DecisionTreeRegressor(max_leaf_nodes=max_leaf_nodes,
                                     random_state=random_state)
        tree.fit(X, y - F)
        trees.append(tree)
        F += lr * tree.predict(xg)
        curves.append(F)
        stage_metrics.append(tree.score(X, y))

    return trees, curves, stage_metrics
```

```

for m in range(n_estimators):
    r = y - F
tree = DecisionTreeRegressor(max_leaf_nodes=max_leaf_nodes,
random_state=rng.randint(0, 10**6))
    tree.fit(X, r)
    trees.append(tree)
    F += lr * tree.predict(X)

Fg = np.full(xg.shape[0], y.mean(), dtype=float)
for t in trees:
    Fg += lr * t.predict(xg)

curves.append(Fg)

mse = float(np.mean((y - F) ** 2))
mae = float(np.mean(np.abs(y - F)))
stage_metrics.append([m+1, mse, mae])

plt.figure(figsize=(6,4))
plt.scatter(X[:,0], y, s=25)
plt.plot(xg[:,0], Fg, linewidth=2, color="red")
plt.title("X vs y")
plt.show()

return F, curves, xg, pd.DataFrame(stage_metrics, columns=["Stage",
"MAE", "MSE"])

```

final_pred, curves, xg, metrics = gradient_boost(X, y, n_estimators=5,
lr=1.0)

```

plt.figure(figsize=(6,4))
plt.scatter(df["X"], df["y"], s=25)
plt.plot(xg[:,0], curves[-1], linewidth=2, color="red")
plt.title("X vs y")
plt.xlabel("X"); plt.ylabel("y")
plt.show()

```

```

plt.figure(figsize=(6,4))
plt.scatter(df["X"], df["y"], s=20)
for c in curves:
    plt.plot(xg[:,0], c, color="red", linewidth=1, alpha=0.8)
plt.title("All Boosting Stages (combined)")

```

```

plt.xlabel("X"); plt.ylabel("y")
plt.show()

plt.figure(figsize=(6,4))
plt.scatter(df["X"], df["y"], s=20)
for c in curves:
    plt.plot(xg[:,0], c, color="red", linewidth=1)
plt.xlabel("X"); plt.ylabel("y")
plt.show()

plt.figure(figsize=(6,4))
plt.scatter(df["X"], df["y"], s=20)
for c in curves:
    plt.plot(xg[:,0], c, color="red", linewidth=1)
plt.xlabel("X"); plt.ylabel("y")
plt.show()

df_out = df.copy()
df_out["y_pred_final"] = final_pred
df_out["residual_final"] = df["y"] - df_out["y_pred_final"]

print("\nTABLE : Predictions & Residuals\n")
print(df_out.head(15))

print("\nTABLE : Stage Metrics (MSE, MAE per boosting stage)\n")
print(metrics)

```

OUTPUT:

